

Alg User Manual

Aleš Bizjak*

Andrej Bauer†

December 9, 2010

Contents

1	Introduction	2
2	Copyright and License	2
3	Installation	3
3.1	Downloading alg	3
3.2	Installation for Linux and MacOS	3
3.2.1	Prerequisites	3
3.2.2	Compiling to native code	4
3.2.3	Compiling to bytecode	4
3.2.4	Installation without Make	4
3.3	Installation for Microsoft Windows	5
4	Input	5
4.1	Comments	5
4.2	General syntactic rules	5
4.3	The Theory keyword	5
4.4	Declaration of operations	6
4.5	Axioms	7
5	Output	7
5.1	Description of the output	7
5.2	Available output formats	8
5.3	Output for further processing	8
6	Command-line Options	8
7	How to use alg efficiently	9

1 Introduction

Alg is a program for enumeration of finite models of algebraic theories. An algebraic theory is given by a signature (a list of constants and operations) and axioms expressed in first-order logic.¹ Examples of algebraic theories include groups, lattices, rings, fields, and many others. Alg can do the following:

*Ales.Bizjak@gmail.com, Faculty of Mathematics and Physics, University of Ljubljana

†Andrej.Bauer@andrej.com, Faculty of Mathematics and Physics, University of Ljubljana

¹Strictly speaking, the axioms of an algebraic theory must be equations, but alg can handle all of first-order logic.

- list or count all non-isomorphic models of a given theory,
- list or count all non-isomorphic indecomposable² models of a given theory.

Currently alg has the following limitations:

- only unary and binary operations are accepted,
- it is assumed that constants denote pairwise distinct elements.

This manual describes how to install and use alg. For a quick start you need Ocaml 3.11 or newer and the menhir parser generator. Compile alg with

```
make
```

and run

```
./alg.native --size 8 theories/unital_commutative_ring.th
```

For usage information type `./alg.native -help` and for examples of theories see the `theories` subdirectory.

Alg is released under the open source simplified BSD License, as detailed in the next section.

2 Copyright and License

Copyright © 2010, Aleš Bizjak and Andrej Bauer

All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
- Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

This software is provided by the copyright holders and contributors “as is” and any express or implied warranties, including, but not limited to, the implied warranties of merchantability and fitness for a particular purpose are disclaimed. In no event shall the copyright holder or contributors be liable for any direct, indirect, incidental, special, exemplary, or consequential damages (including, but not limited to, procurement of substitute goods or services; loss of use, data, or profits; or business interruption) however caused and on any theory of liability, whether in contract, strict liability, or tort (including negligence or otherwise) arising in any way out of the use of this software, even if advised of the possibility of such damage.

3 Installation

3.1 Downloading alg

Alg is available at <http://hg.andrej.com/alg/>. You have three options:

1. download the ZIP file with source code from

<http://hg.andrej.com/alg/archive/tip.zip>

²A model is indecomposable if it cannot be written as a non-trivial product of two smaller models.

2. clone the repository with the Mercurial revision control system:

```
hg clone http://hg.andrej.com/alg/
```

3. download a precompiled executable for your architecture from

```
http://hg.andrej.com/alg/file/tip/precompiled
```

if one is available. If you choose this option, make sure that you still obtain the ZIP file because the `theories` subdirectory contains a number of useful examples.

3.2 Installation for Linux and MacOS

3.2.1 Prerequisites

To compile `alg` you need the `Make` utility, Ocaml 3.11 or newer, and the `menhir` parser generator higher. We will assume you have `Make`. You can get Ocaml and `menhir` in several ways:

1. On Ubuntu, install the packages `ocaml` and `menhir`:

```
sudo apt-get install ocaml menhir
```

Similar solutions are available on other Linux distributions.

2. On MacOS the easiest way to install Ocaml and `menhir` is with the `macports` utility:

```
sudo port install ocaml
sudo port install caml-menhir
```

3. If you have `GODI` installed then you already have Ocaml. Install `menhir` with the `godi_console` command, if you do not have it yet.

4. Ocaml is also available from

```
http://caml.inria.fr/
```

and `menhir` from

```
http://pauillac.inria.fr/~fpottier/menhir/
```

3.2.2 Compiling to native code

To compile `alg`, type `make` at the command line. If all goes well `ocamlbuild` will generate a subdirectory `_build` and in it the `alg.native` executable. It will also create a link to `_build/alg.native` from the top directory. To test `alg` type

```
./alg.native --count --size 8 theories/group.th
```

It should tell you within seconds that there are 5 groups of size 8.

We provided only a very rudimentary installation procedure for `alg`. First edit the `INSTALL_DIR` setting in `Makefile` to set the directory in which `alg` should be installed, then run

```
sudo make install
```

This will simply copy `_build/alg.native` to `$(INSTALL_DIR)/alg`. You may also wish to stash the `theories` subdirectory somewhere for future reference.

3.2.3 Compiling to bytecode

If your version of Ocaml does not compile to native code you can try compiling to bytecode with

```
make byte
```

This will generate a (significantly slower) `alg.byte` executable.

3.2.4 Installation without Make

If you do not have the Make utility (how can that be?) you can compile `alg` directly with `ocamlbuild`:

```
ocamlbuild -use-menhir alg.native
```

To install `alg` just copy `_build/alg.native` to `/usr/local/bin/alg` or some other reasonable place.

3.3 Installation for Microsoft Windows

Sorry, this has not been written yet. But if you have Make and Ocaml 3.11 and `menhir`, you should be able to just follow the instructions for Linux.

Note that a Windows precompiled executable may be available at

<http://hg.andrej.com/alg/tip/precompiled/>

4 Input

An `alg` input file has extension `.th` and it describes an algebraic theory. The syntax vaguely follows the syntax of the Coq proof assistant. A typical input file might look like this:

```
# The axioms of a group.
Theory group.
Constant 1.
Unary inv.
Binary *.
Axiom unit_left: 1 * x = x.
Axiom unit_right: x * 1 = x.
Axiom inverse_left: x * inv(x) = 1.
Axiom inverse_right: inv(x) * x = 1.
Axiom associativity: (x * y) * z = x * (y * z).
```

There is an optional **Theory** declaration which names the theory, then we have declarations of constants, unary and binary operations, and after that there are the axioms. The precise syntax rules are as follows.

4.1 Comments

Comments are written as in Python, i.e., a comment begins with the `#` symbol and includes everything up to the end of line.

4.2 General syntactic rules

An `alg` input file consists of a sequence of declarations (**Theory**, **Constant**, **Unary**, **Binary**) and axioms (**Axiom**, **Theorem**). Each declaration and axiom is terminated with a period.

4.3 The Theory keyword

You may give a name to your theory with the declaration

```
Theory theory_name.
```

at the beginning of the input file, possibly preceded by comments and whitespace. The theory name consists of letters, numbers, and the underscore. If you do not provide a theory name, alg will deduce one from the file name.

4.4 Declaration of operations

The declarations

```
Constant c1 c2 ... ck .
Unary u1 u2 ... um .
Binary b1 b2 ... bn .
```

are used to declare constants, unary, and binary operations respectively. You may declare several constants or operations with a single declaration, or one at a time. You may mix declarations and axioms, although it is probably a good idea to declare the constants and operations first.

A constant may be any string of letters, digits and the underscore character. In particular, a constant may consist just of digits, for example 0 or 1.

Unary and binary operations may be strings of letters, digits and the underscore character. For example, if we declare

```
Unary inv .
Binary mult .
```

then we can write expressions like `mult(x, inv(y))`. It is even possible to declare operations whose names are strings of digits, for example:

```
Unary 3 ten .
Binary + .
Axiom: 3(3(x)) + x = ten(x) .
```

Alternatively, we can use *infix* and *prefix* operators. These follow the Ocaml rules for infix and prefix notation. An operator is a string of symbols

`! $ % & * + - / \ : < = > ? @ \ ^ | ~`

where:

- a *prefix operator* is one that starts with `?`, `!` or `~`. It can be used as a unary operation.
- *infix operators* can be used as binary operations and have four levels of precedence, listed from lowest to highest:
 - left-associative operators starting with `|`, `&`, `$`
 - right-associative operators starting with `@` and `^`
 - left-associative operators starting with `+`, `-`, and `\`
 - left-associative operators starting with `*`, `/`, and `%`
 - right-associative operators starting with `**`.

An operator \circ is *left-associative* if $x \circ y \circ z$ is understood as $(x \circ y) \circ z$, and *right-associative* if $x \circ y \circ z$ is understood as $x \circ (y \circ z)$. If you look at the above list again, you will notice that operators have the expected precedence and associativity. However, if you are unsure about precedence, it is best to use a couple of extra parentheses.

4.5 Axioms

An axiom has the form

`Axiom [name]: <formula>.`

or

`Theorem [name]: <formula>.`

There is no difference between an axiom and a theorem as far as alg is concerned. We use **Axiom** for the actual axioms and **Theorem** for statements that are consequences of axioms and are worth including in the theory because they make alg run faster, see Section 7.

The optional *[name]* is a string of letters, digits and the underscore characters. The *<formula>* is a first-order formula built from the following logical operations, listed in order of increasing precedence:

$\forall x. \phi$	is written as	<code>forall x, ϕ,</code>
$\exists x. \phi$	is written as	<code>exists x, ϕ,</code>
$\phi \Leftrightarrow \psi$	is written as	<code>$\phi \leftrightarrow \psi$ or $\phi \Leftrightarrow \psi$,</code>
$\phi \Rightarrow \psi$	is written as	<code>$\phi \rightarrow \psi$ or $\phi \Rightarrow \psi$,</code>
$\phi \vee \psi$	is written as	<code>$\phi \vee \psi$ or ϕ or ψ,</code>
$\phi \wedge \psi$	is written as	<code>$\phi \wedge \psi$ or ϕ and ψ,</code>
$\neg \phi$	is written as	<code>not ϕ,</code>
$s = t$	is written as	<code>$s = t$,</code>
$s \neq t$	is written as	<code>$s <> t$ or $s \neq t$,</code>
\top and \perp	are written as	<code>True and False, respectively.</code>

An iterated quantification $\forall x_1. \forall x_2. \dots \forall x_n. \phi$ may be written as

`forall x_1 x_2 ... x_n , ϕ .`

and similarly for \exists .

Axioms may contain free variables. Thus we can write just

`Axiom: $x + y = y + x$.`

instead of

`Axiom: forall x y, $x + y = y + x$.`

5 Output

5.1 Description of the output

The output of alg is meant to be self-explanatory. Nevertheless, here is what the output consists of:

Title: the name of the theory.

Theory: the input file, which can be suppressed with `--no-source` command-line option.

Models: a list of the models found. This can be suppressed with the `--count` command-line option. Each model has a name *theory_name_n_m* where *n* is the model size and *m* is the model sequence number. If a model can be decomposed, a decomposition into indecomposable factors is given.³ Tables of all the operations are displayed.

Counts: a table showing how many models of each size were found. If more than three sizes were considered, alg also provides a URL to query the counts at <http://oeis.org/>, the On-Line Encyclopedia of Integer Sequences.

³Please note that in general such a decomposition is *not* unique.

5.2 Available output formats

Alg supports several output formats. The default output format is plain text and it is sent to the screen. You can choose a different format with the `--format` command-line option, and you can send the output to a file with the `--output` option. If you specify an output file but no format, alg guesses the correct format from the output filename. To see which formats are supported by alg, type `alg --help`.

5.3 Output for further processing

Alg can output the models in JSON format, which is suitable for further processing. Select the JSON format by specifying the command-line option `--format json` or an output file with the `.json` extension.

In Python you can import JSON data from a file `mystuff.json` like this:

```
import json
with open('mystuff.json', 'r') as f:
    mystuff = json.load(f)
```

In Mathematica just use the `Import["mystuff.json"]` command. Most other common programming languages have support for processing JSON data.

6 Command-line Options

Alg is used as

```
alg --size <sizes> [options] <theory.th>
```

where *theory.th* is the input file, and the options are:

`--size <sizes>` A comma-separated list of sizes that alg should consider. You can also specify a size interval of the form `m-n`. For example, `1,2,5-8` would mean that we consider sizes 1, 2, 5, 6, 7, 8.

`--count` Do not print out the models, just report the counts.

`--format <format>` Output in the given format. Supported formats are `text`, `html`, and `latex`.

`--indecomposable` Output only indecomposable models, i.e., those that are not products of smaller models.

`--no-products` Do not try to generate models as products of smaller models. Use this option if you know that a model cannot be a product of smaller ones. For example, a field can never be a product of two fields. If all of your axioms are equations, then you should *not* use this option.

`--no-source` Do not include the theory in the output.

`--output <filename>` Output to the given file rather than to screen.

`--help` Print help.

7 How to use alg efficiently

You should always keep in mind the fact that alg performs a brute force search with a few optimizations. In the worst case its running time is doubly exponential in the size of the models because there are n^{n^2} tables for a binary operation on a set of size n .

Alg is optimized for *equational* theories, i.e., those whose axioms are equations (semigroups, monoids, groups, rings, lattices, etc., but *not* integral domains and fields). Alg takes advantage of commutativity, associativity and idempotent laws, and to a smaller extent of other kinds of equational laws, such as absorption and distributivity.

Alg checks axioms which are not equations, but does not perform any optimizations based on them. You should have as few non-equational axioms as possible. Furthermore, you should always push all the quantifiers inside. For example, instead of

```
Axiom: forall x, exists y, x <> 0 -> x * y = 1.
```

you should write

```
Axiom: forall x, x <> 0 -> (exists y, x * y = 1).
```

The best kind of axioms are those that allow alg to immediately fill in a whole column or row. Typically these are axioms about neutral elements, such as $1 \cdot x = x$ and $0 + x = x$. As a rule of thumb, every such axiom will increase the maximum manageable size by one.

In general alg performs better if it is given more axioms and theorems, because each additional statement cuts down the possibilities. Thus you *should* include theorems which already follow from other axioms. For example:

- state *both* $1 \cdot x = x$ and $x \cdot 1 = x$, even if one of them follows from the other,
- more generally, state all versions of a symmetric equation, even if they all follow from one of them,
- state laws like $0 \cdot x = 0$ (and also $x \cdot 0 = 0$), even if they follow from other axioms.

You should declare as many constants and as few operations as possible. A typical example is the theory of lattices. For *finite* structures the following are equivalent theories:

- a lattice with operations \wedge, \vee ,
- a bounded lattice with operations \wedge, \vee and constants 0, 1,
- a \vee -semilattice with operation \vee and constant 0,
- a bounded \vee -semilattice with operation \vee and constants 0, 1.

The best choice is the last one because it has just fewest operation and most constants. Indeed, figuring out that there are 53 lattices of size 7 takes 250 times longer with the theory of a lattice than with the theory of a bounded \vee -semilattice.

Lastly, we should mention that alg generates tables in the order in which the operations are declared. Sometimes it is much easier to generate tables for one operation than another, so you should experiment by switching the order of declarations. For example, the theory of a ring works much faster if addition $+$ is declared before multiplication \times .